

8.65

Stochastic Processes: Assignment 1

Group 1: Javier Esteban Aragoneses, Mauricio Marcos Fajgenbaum, Danyu Zhang, Daniel Alonso

November 27th, 2020

Do not include scripts unless
on the appendix

Importing libraries

```
library(markovchain)  
library(matlib)
```

→ See instructions

Function to solve the problems

```
# Thanks prof!  
matrixpower <- function(M,k) {  
  if(dim(M)[1] != dim(M)[2]) return(print("Error: matrix M is not square"))  
  if (k == 0) return(diag(dim(M)[1]))  
  if (k == 1) return(M)  
  if (k > 1) return(M %*% matrixpower(M, k-1))  
}
```

Problem 1

a) 10.5

Markov chain criteria:

1- The probability of being in a state only depends on the previous state.

2- It's a stochastic process.

$X =$ The chain hits state j at time n

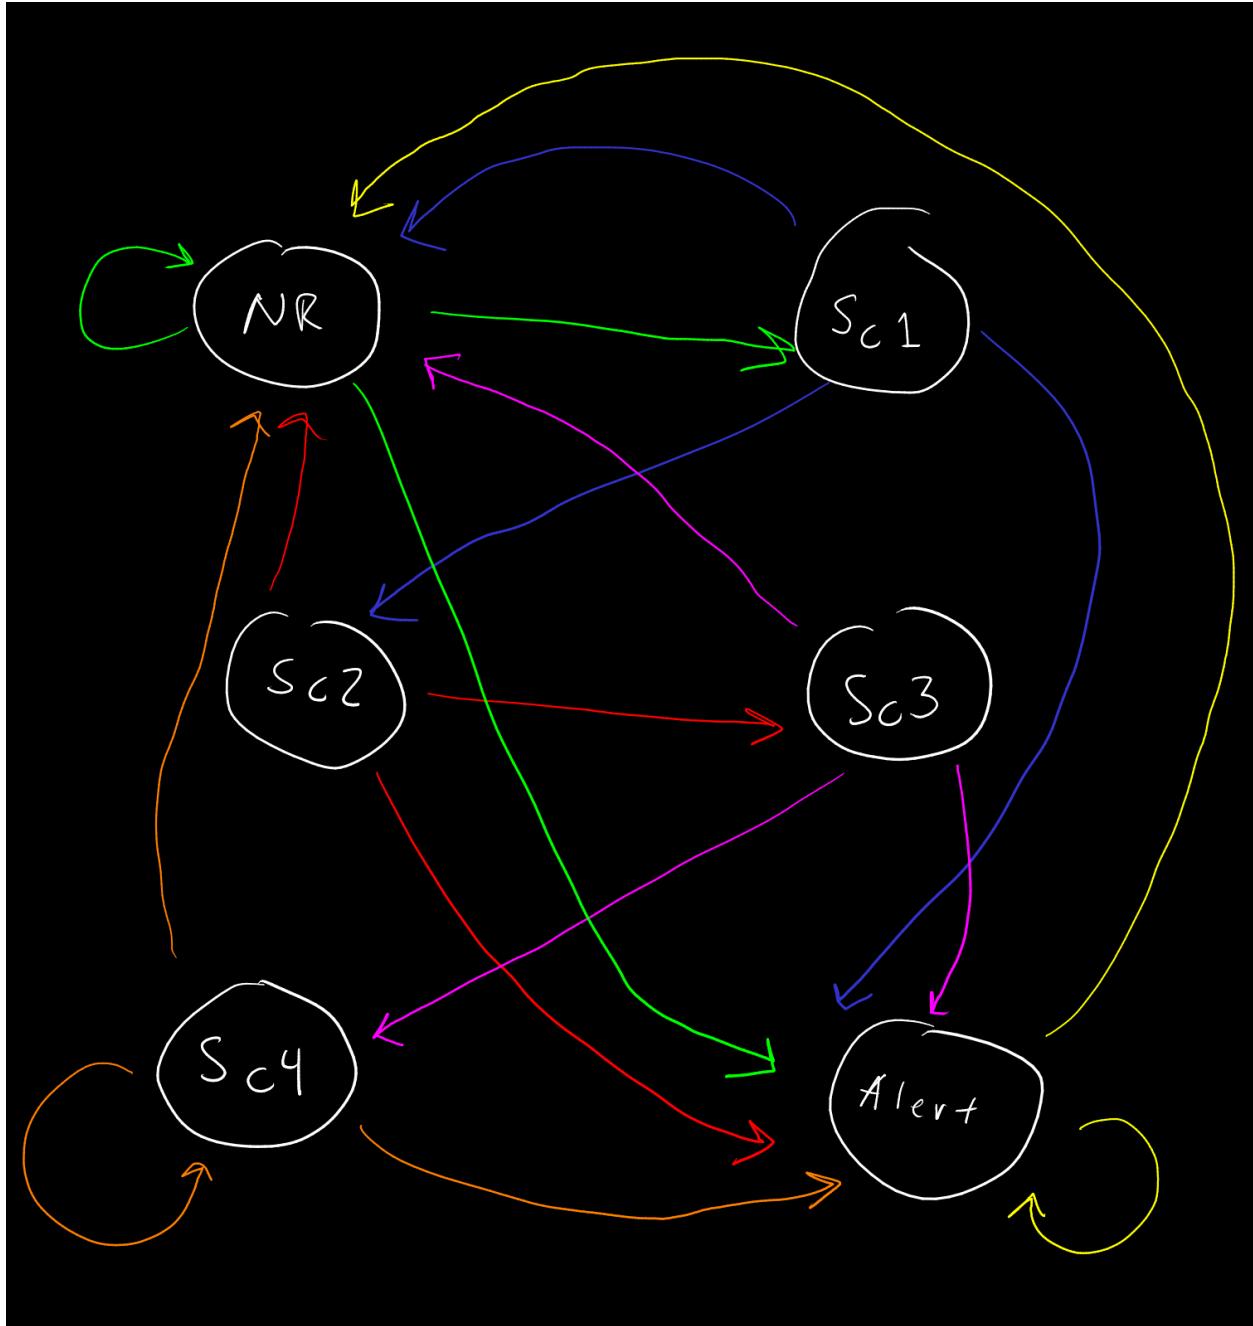
X_n is the scenario at time n

All states have finite expected return times and there is only a single communication class so the MC is irreducible, therefore its stationary distribution is **unique**.

You need to explain why the Markov property holds in this case

But why is that?

Ok for the states and
possible transitions 0.5



Transition
Figure 1: Graph for prob.1

1

Again, you should explain what you do, with your own words, formulae and then, if necessary, refer to the code in the appendix.

b) We have first calculated the relative frequencies manually.

```
# importing the data
load('PollutionMadrid.RData')
data <- X[1,] # selecting the first line (as we are group 1)
mat <- matrix(rep(0,36), nrow=6, byrow=T) # creating an empty 6x6 matrix
# we substitute the characters with numbers for easier manipulation
for (i in 1:length(data)) {
  if (data[i] == "Alert") {
    data[i] = 1
  } else if (data[i] == "NR") {
    data[i] = 2
  } else if (data[i] == "Sc1") {
    data[i] = 3
  } else if (data[i] == "Sc2") {
    data[i] = 4
  } else if (data[i] == "Sc3") {
    data[i] = 5
  } else if (data[i] == "Sc4") {
    data[i] = 6
  }
}

# Count transitions and add them to the matrix
data <- as.numeric(data)
for (i in 1:length(data)) {
  mat[data[i],data[i+1]] = mat[data[i],data[i+1]] + 1
}
# matrix wraps around so we count the last element > first element transition
mat[data[1460],data[1]] = mat[data[1460],data[1]] + 1

# Frequency table in order to create relative frequencies in transition matrix
tbl <- table(data)
for (i in 1:length(tbl)) {
  mat[i,] = mat[i,]/tbl[i]
}
mat
#>      [,1]     [,2]     [,3]     [,4]     [,5]     [,6]
#> [1,] 0.00000000 1.0000000 0.00000000 0.0000000 0.0000000 0.0000000
#> [2,] 0.00000000 0.9529851 0.04701493 0.0000000 0.0000000 0.0000000
#> [3,] 0.01587302 0.4920635 0.00000000 0.4920635 0.0000000 0.0000000
#> [4,] 0.00000000 0.5806452 0.00000000 0.0000000 0.4193548 0.0000000
#> [5,] 0.23076923 0.3846154 0.00000000 0.0000000 0.0000000 0.3846154
#> [6,] 0.00000000 0.5555556 0.00000000 0.0000000 0.0000000 0.4444444
```

✓

0.75

Ok, but indicate what the order of states is (explicitly, not to have to go and look for it in the code)

We then tested using the *markovchain* package in order to confirm our results.

```
data <- X[,]  
markovchainFit(data)$estimate  
#> MLE Fit  
#> A 6 - dimensional discrete Markov Chain defined by the following states:  
#> Alert, NR, Sc1, Sc2, Sc3, Sc4  
#> The transition matrix (by rows) is defined as follows:  
#>          Alert      NR      Sc1  Sc2      Sc3      Sc4  
#> Alert 0.00000000 1.0000000 0.00000000 0.0 0.0000000 0.0000000  
#> NR    0.00000000 0.9529851 0.04701493 0.0 0.0000000 0.0000000  
#> Sc1   0.01612903 0.4838710 0.00000000 0.5 0.0000000 0.0000000  
#> Sc2   0.00000000 0.5806452 0.00000000 0.0 0.4193548 0.0000000  
#> Sc3   0.23076923 0.3846154 0.00000000 0.0 0.0000000 0.3846154  
#> Sc4   0.00000000 0.5555556 0.00000000 0.0 0.0000000 0.4444444
```



What can you say of the comparison of your estimates and the possible transitions between states that you had argued in part a

According to our probabilities shown in the graph. There are 3 arrows with probability 0. This is due to the fact that in the data there are zero transitions from $Sc2 \rightarrow Alert$, $Sc4 \rightarrow Alert$, $NR \rightarrow Alert$, $Alert \rightarrow Alert$.

This is logical given that it is very unlikely to hit an alert state. Unlike the rest of the states.

Later it will be shown that there's a unique stationary distribution (see 1d).

0.25

although possible,
that is why you have
an arrow in your
transition graph



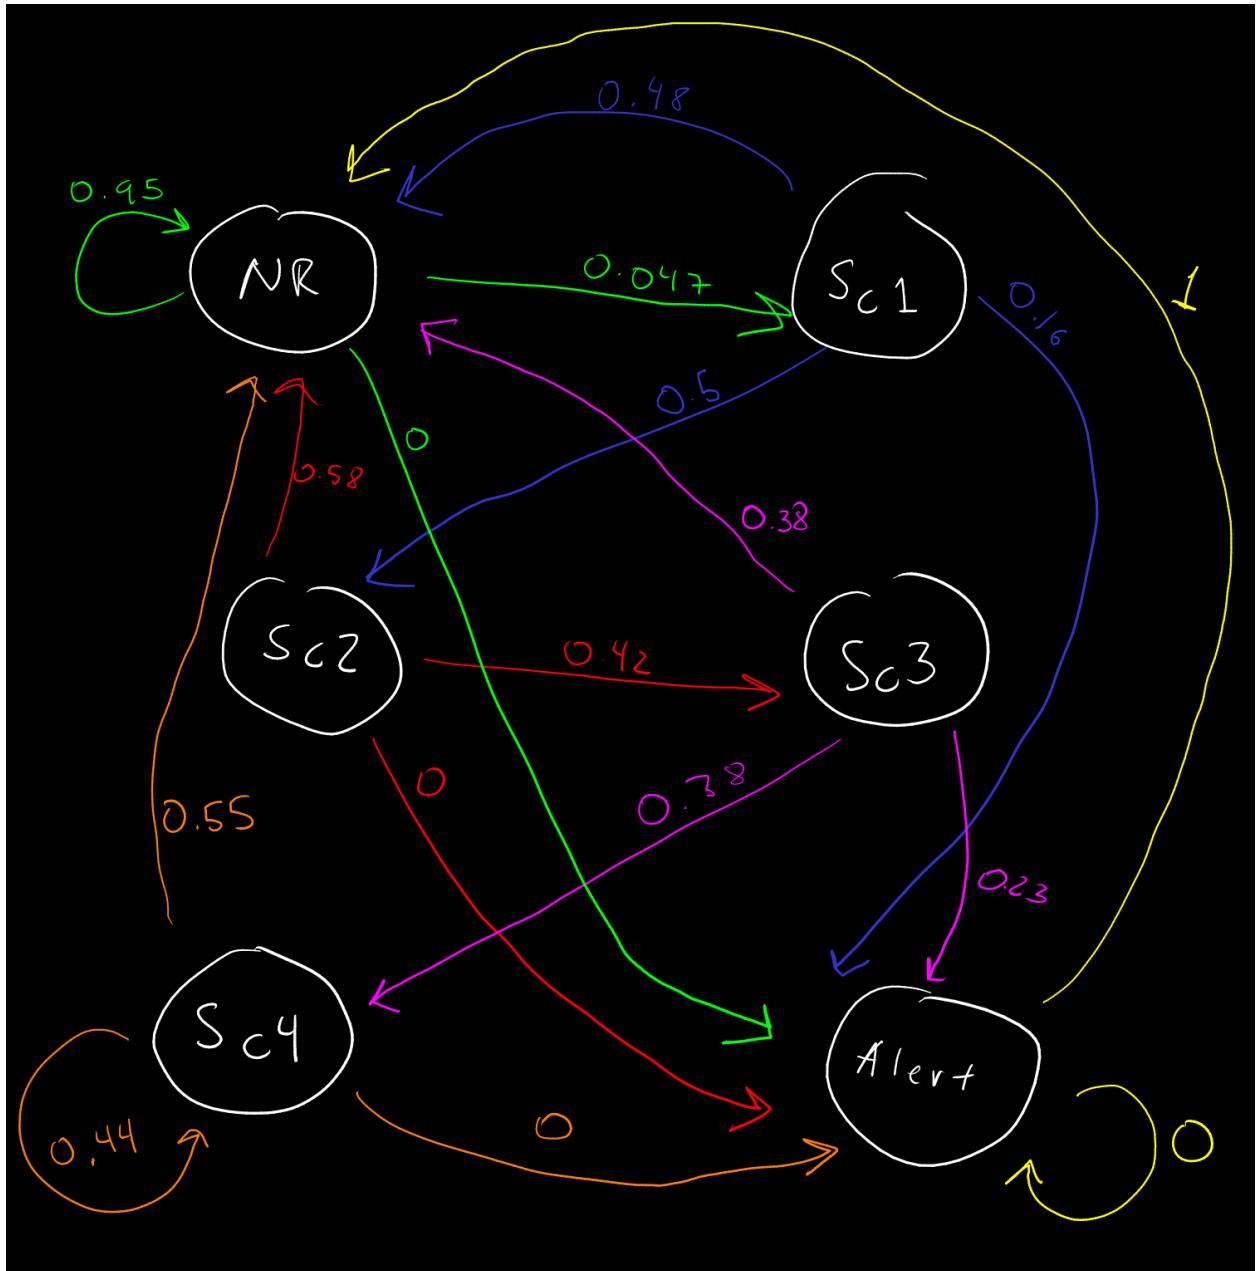


Figure 2: Graph with probabilities for problem 1 (b)

estimated

c) 0.71

Given that the first state of the chain is NR. We see the following 7 states:

```
data[1:7]
#> [1] "NR" "NR" "NR" "NR" "NR" "NR" "NR"
```

And we calculate the probability as follows:

```
mat[2,2]^7
#> [1] 0.713843
```

We can see the probability is 0.713843

d) 1.5

We can see that because we have a unique solution to the system, we have a unique stationary distribution.

function to calculate stationary distribution of a transition matrix

```
stationary_dist <- function(P) {
  dim = sqrt(length(P))
  A = P - diag(dim)
  b = matrix(c(1,rep(0,dim-1)),nrow=dim,byrow=T)
  A[,1] <- rep(1,dim)
  print("The solution is the following:")
  return(matlib::Solve(t(A), b))
}

stat_dist <- stationary_dist(mat)
#> [1] "The solution is the following:
#>   x1          =  0.00273973
#>   x2          =  0.91780822
#>   x3          =  0.04315068
#>   x4          =  0.02123288
#>   x5          =  0.00890411
#>   x6          =  0.00616438
stat_dist
#> [1] "x1          =  0.00273973" " x2          =  0.91780822"
#> [3] "   x3          =  0.04315068" "   x4          =  0.02123288"
#> [5] "   x5          =  0.00890411" "   x6          =  0.00616438"
```

Which yields the following stationary distribution:

$$\pi_1 = 0.00273973$$

$$\pi_2 = 0.91780822$$

$$\pi_3 = 0.04315068$$

$$\pi_4 = 0.02123288$$

$$\pi_5 = 0.00890411$$

$$\pi_6 = 0.00616438$$

0.71
That would be $P(X_f = NR | X_0 = NR)$
for instance

Ok, but what is that you are calculating? How do you get to P_{22}^7 ? You have to explain things!

What system?
Say it explicitly!

0.75

Again, reminding what state corresponds to each number 1 to 6 would be useful.

what proportions? explain
why you can do that

Comparing with the proportions we get from our data:

```
rel_error = c()
props = table(data)/length(data)
results <- c(0.00273973, 0.91780822, 0.04315068, 0.02123288, 0.00890411, 0.00616438)
for (i in 1:length(props)) {
  rel_error[i] <- abs(props[i]-results[i])/results[i]
}
rel_error
#> [1] 1.449998e-06 8.955223e-10 1.142857e-07 1.548387e-07 4.615384e-08
#> [6] 5.777781e-07
```

0.5

We can see our relative errors are all quite low ($< 1 \times 10^{-5}$) so we could say that the difference between both estimations is nearly negligible.

e) 0.7

Proof that this MC has a limiting distribution

but what is the conceptual difference between the 2 approaches?

We have previously proven that this MC is irreducible. Now we want to prove that this MC is also aperiodic.

We can easily prove this by checking the paths from NR to the nodes it's connected with:

1st path: 3 steps: $NR \rightarrow Sc1 \rightarrow Sc2 \rightarrow NR$

2nd path: 2 steps: $NR \rightarrow Sc1 \rightarrow NR$

...

Because the $\{3, 2, \dots\}$ and 2 elements of this set are prime numbers, the greatest common divisor of this set is always going to be 1.

This applies for all the other nodes as they all belong to the same class.

~~Because it is finite!~~ Therefore this chain has a unique stationary distribution which coincides with its limiting distribution.

0.5

What does this mean, in terms of pollution episodes?

The proportion of pollution episodes depends on our limiting distribution, i.e. the long run proportion of pollution episodes with Scenario 1 will be the value of $Sc1$ in our limiting distribution. If we look ahead in the future, and we want to know the probability of being in a specific pollution episode scenario, p_{ij} is the probability of being in scenario j .

Taking the a high power of our transition matrix we get the following limiting distribution:

```
mp <- matrixpower(mat, 120)
mp
#>      [,1]     [,2]     [,3]     [,4]     [,5]     [,6]
#> [1,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
#> [2,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
#> [3,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
#> [4,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
#> [5,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
#> [6,] 0.002739726 0.91780822 0.04315068 0.02123288 0.00890411 0.006164384
```

0.2

Ok, but this is also true if a stationary (not limiting) distribution exist.

f) 0.5

If we take the sum of the probability of scenarios 3, 4 and Alert: $P(Sc3) + P(Sc4) + P(Alert)$ multiplied by the amount of days in a year, we get the amount of days in a year with such driving restriction:

365*(mp[1,1]+mp[1,5]+mp[1,6])
#> [1] 6.5



0.5

We get that the amount of days where driving is forbidden is 6.5 days. Approximately a week.

3.75

Problem 2

a)

1.2

We set up the following system of equations:

$$\begin{cases} \sum_{i \in \mathbb{N} \cup \{0\}} \pi_i P_{i,0} = \pi_1 \\ \sum_{i \in \mathbb{N} \cup \{0\}} \pi_i = 1 \\ (1-p)\pi_1 = \pi_2 \\ \dots \\ (1-p)\pi_{n-2} = \pi_{n-1} \\ \dots \end{cases}$$

ok but that should be π_0 to be consistent with notation for P

For the first equation, each $P_{i,0} = p$, therefore:

$$\sum_{i \in \mathbb{N} \cup \{0\}} P_{i,0} \pi_i = \pi_1 \Rightarrow p \sum_{i \in \mathbb{N} \cup \{0\}} \pi_i = \pi_1$$

And we get:

$$p = \pi_1$$

For the rest of the equations we have the following:

$$(1-p)p = \pi_2, (1-p)^2p = \pi_3, \dots, (1-p)^{n-1}p = \pi_n, \dots$$

Then, we get the following result:

$$p + (1-p)p + (1-p)^2p + \dots + (1-p)^{n-1}p + \dots = \sum_{i \in \mathbb{N} \cup \{0\}} \pi_i = 1$$

The left side of this equation can be summarized by a summation:

$$p \sum_{i \in \mathbb{N} \cup \{0\}} (1-p)^k = \sum_{i \in \mathbb{N} \cup \{0\}} \pi_i = 1$$

The summation on the left hand side is a geometric series where $|(1-p)| < 1$, therefore:

$$p * \frac{1}{1-(1-p)} = 1$$

This series confirms that $\pi P = \pi$ because:

$$[p \quad (1-p)p \quad \dots \quad (1-p)^n p \quad \dots] \begin{bmatrix} p & 1-p & \dots & 0 \\ \vdots & \vdots & \ddots & \end{bmatrix} = \pi$$

Performing this matrix product we get:

$$\begin{cases} \pi_1 = p \\ \pi_2 = p(1-p) \\ \pi_3 = p(1-p)^2 \\ \dots \\ \pi_n = p(1-p)^{n-1} \\ \dots \end{cases} \quad \begin{array}{l} \text{TO} \\ \text{PI} \\ \text{TII} \\ , \\ \text{Tn-1} \\ , \end{array}$$

Basically:

$$\pi_i = p(1-p)^{i-1} \text{ for all } i \in \mathbb{N} \cup \{0\}$$

Which is the stationary distribution of the chain.

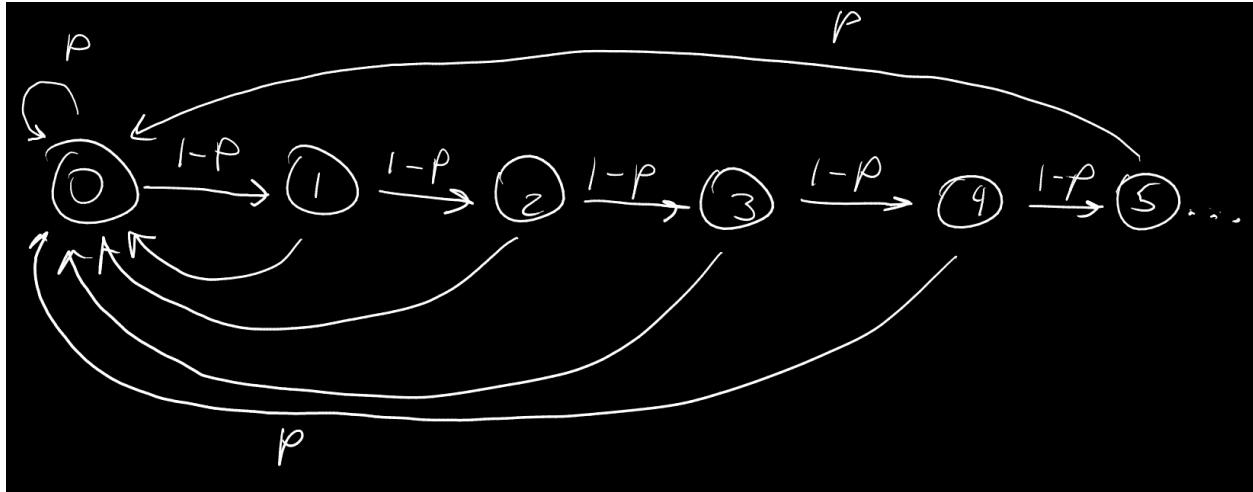


Figure 3: Graph for problem 2a

Because our MC is an irreducible infinite state MC, state 0 is a recurrent state and given that the matrix is irreducible, it is a recurrent MC.

- 1 - We have a single communication class therefore it's irreducible.
- 2 - Zero has period 1, so all states have period 1 and therefore the MC aperiodic.
- 3 - We calculate the probability of returning to 0 in exactly n steps:

$$\begin{aligned} f_{00} &= \sum_{n>1} P(\text{returning to 0 in exactly } n \text{ steps}) \\ &= \sum_{n>1} (f_{00})^n = p + (1-p)p + \dots + (1-p)^n p + \dots \\ &= p \sum_{n>1} (1-p)^k = \frac{p}{1-(1-p)} = \frac{p}{p} = 1 \end{aligned}$$

Then we calculate the expected return time to zero:

$$\begin{aligned} E[\text{return time to 0}] &= \sum_{n>0} n(f_{00})^n = 1 * p + 2p(1-p) + \dots + n(1-p)^{n-1} p + \dots \\ &= p \sum_{n>0} (1-p)^{n-1} = p \sum_{n>0} \frac{n(1-p)^n}{(1-p)} \\ &= \frac{p}{1-p} \sum_{n>0} n(1-p)^n \text{ with } p \neq 1, 0 \end{aligned}$$

We know that:

$$n(1-p)^n = \frac{n(1-p)^n}{(n-1)(1-p)^{n-1}}$$

and using the ratio test for this series:

$$\lim_{n \rightarrow \infty} \frac{|n(1-p)^n|}{|(n-1)(1-p)^{n-1}|} = \lim_{n \rightarrow \infty} \frac{|n|}{|n-1|} (1-p) = 1-p < 1$$

Therefore our series converges for all $p \neq 0$.

State 0 is positive recurrent and so are all the rest of the states in the MC.

As we have a unique stationary distribution π , $\pi_i = \frac{1}{\mu_i}$ and all states have expected finite return times then we have:

$$E[T_i | X_0 = i] = \mu_i = \frac{1}{\pi_i} \quad \text{And specifically if } i = 0?$$

cause you have a little

Since it is an irreducible aperiodic MC for which a stationary dist. exists \Rightarrow it is unique and $\pi_i = \frac{1}{E[T_i | X_0 = i]}$

Not necessary

b) 0.75 because you have discussed
in a)

As it is aperiodic and irreducible markov chain, it has a limiting distribution, as all states are positive recurrent.

c) 0.9

We define the following function to simulate n steps with initial value p .

```
sim <-function(x0,n,p){
  n <- n-2
  x = rep(0,n)
  for (i in 1:n) {
    u=runif(1)
    if (u <= p) {
      x[i+1] = x[i] + 1
    } else {
      x[i+1] = 0
    }
  }
  return(c(x0,x))
}
```

$$\begin{aligned} \rightarrow \text{ok, but then} \\ P_{i,i+1} &= p \\ P_{i,0} &= 1-p \end{aligned}$$

We simulate $n = 30$ steps to test the function with initial value $x_0 = 100$ and $p = 0.95$:

```
sim(100,30,0.95)
#> [1] 100 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
#> [20] 18 19 20 21 22 23 24 25 26 27 28
```

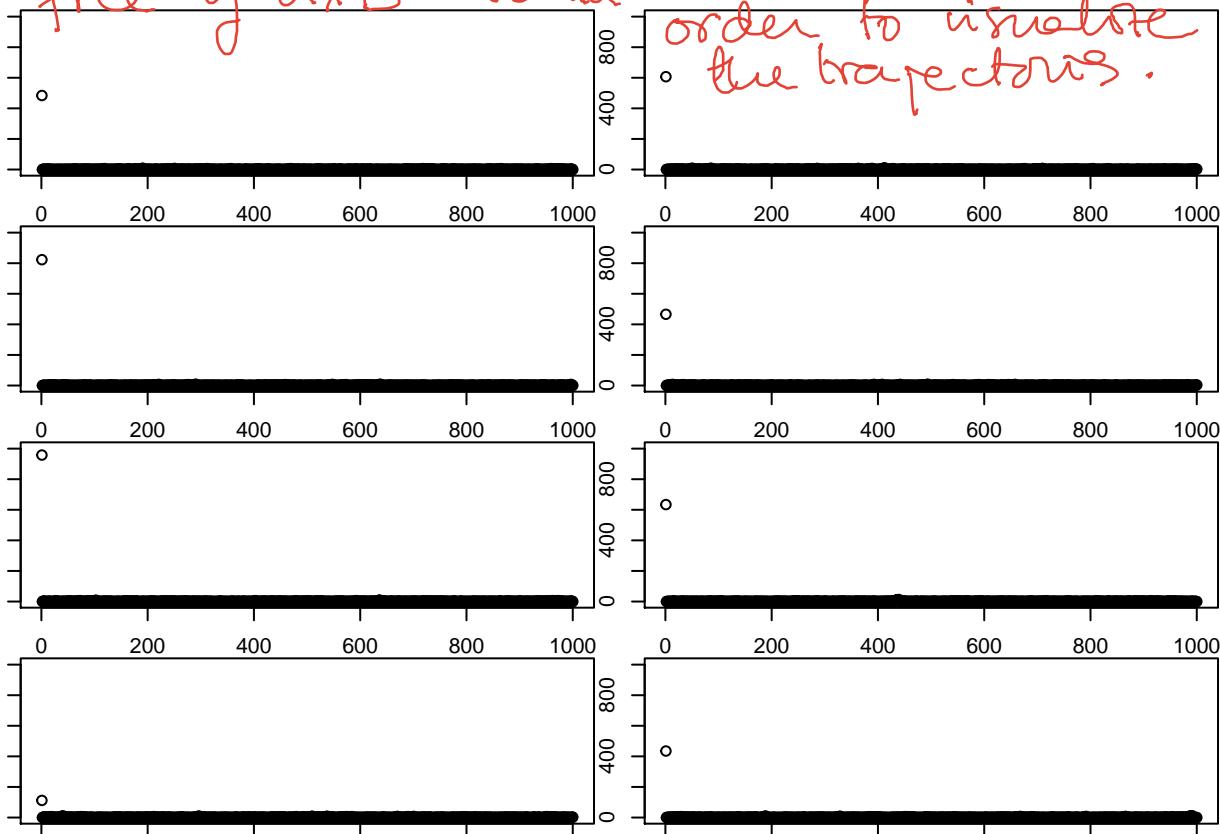
d) 0.9

We generate 8 trajectories of length 1000 with different intial values between [100,1000]

```
# simulating trajectories
initial_x = sample((100:1000), 8)
trajs = matrix(rep(0,8*1000), nrow=1000, byrow=T)
for (i in 1:8) {
  trajs[,i] <- sim(initial_x[i], 1000, 1/2)
}

# plotting trajectories
par(mfrow=c(4,2), mar=c(1,1,1,1))
for (i in 1:8) {
  plot(trajs[,i], ylim=c(0,1001))
```

Lines instead of bullets and free y-axis would have been better in order to visualize the trajectories.



The limiting distribution of this chain does not really depend at all on the first value. As such initial value could be a high number (between 100-1000), given that the probability corresponding to our group number ($k = 1$) produces a very high $p = \frac{1}{1+k} = \frac{1}{2}$, therefore the chain drops very quickly to 0 and takes a really large number of steps to go back to its initial value.

Our maximums per simulation are always the initial value:

```
maximums = c()
for (i in 1:8) {
  maximums[i] = max(trajs[,i])
}
maximums
#> [1] 484 607 823 466 958 634 112 435
```

it doesn't go back
it stays on its
Stationary
distribution
Geom(p)

While the second highest values will be the followings:

```
maximums = c()
for (i in 1:8) {
  maximums[i] = max(trajs[,i][2:1000])
}
maximums
#> [1] 8 10 8 8 9 13 10 12
```

Which are all significantly lower than our initial values.